

# Learning MNIST with Neural Networks with backprop library

Justin Le

The *backprop* library performs back-propagation over a *heterogeneous* system of relationships. back-propagation is done automatically (as reverse-mode automatic differentiation), and you work with your values as if you were writing normal functions with them, with the help of lens.

Repository source is on github, and docs are on hackage.

If you're reading this as a literate haskell file, you should know that a rendered pdf version is available on github.. If you are reading this as a pdf file, you should know that a literate haskell version that you can run is also available on github!

```
{-# LANGUAGE BangPatterns           #-}
{-# LANGUAGE DataKinds             #-}
{-# LANGUAGE DeriveGeneric         #-}
{-# LANGUAGE FlexibleContexts       #-}
{-# LANGUAGE GADTs                 #-}
{-# LANGUAGE LambdaCase            #-}
{-# LANGUAGE ScopedTypeVariables   #-}
{-# LANGUAGE TemplateHaskell       #-}
{-# LANGUAGE TupleSections         #-}
{-# LANGUAGE TypeApplications      #-}
{-# LANGUAGE ViewPatterns          #-}
{-# OPTIONS_GHC -fno-warn-incomplete-patterns #-}
{-# OPTIONS_GHC -fno-warn-orphans      #-}
{-# OPTIONS_GHC -fno-warn-unused-top-binds  #-}

import           Control.DeepSeq
import           Control.Exception
import           Control.Lens hiding ((<.>))
import           Control.Monad
import           Control.Monad.IO.Class
import           Control.Monad.Trans.Maybe
import           Control.Monad.Trans.State
import           Data.Reflection
import           Data.Bitraversable
import           Data.Foldable
import           Data.IDX
import           Data.List.Split
import           Data.Time.Clock
import           Data.Traversable
import           Data.Tuple
import           GHC.Generics           (Generic)
import           GHC.TypeLits
import           Numeric.Backprop
import           Numeric.LinearAlgebra.Static hiding (dot)
import           Text.Printf
```

```

import qualified Data.Vector           as V
import qualified Data.Vector.Generic   as VG
import qualified Data.Vector.Unboxed   as VU
import qualified Numeric.LinearAlgebra as HM
import qualified System.Random.MWC     as MWC
import qualified System.Random.MWC.Distributions as MWC

```

## Types

For the most part, we're going to be using the great *hmatrix* library and its vector and matrix types. It offers a type `L m n` for  $m \times n$  matrices, and a type `R n` for an  $n$  vector.

First things first: let's define our neural networks as simple containers of parameters (weight matrices and bias vectors).

First, a type for layers:

```

data Layer i o =
  Layer { _lWeights :: !(L o i)
        , _lBiases  :: !(R o)
        }
  deriving (Show, Generic)

```

```

instance NFData (Layer i o)
makeLenses ''Layer

```

And a type for a simple feed-forward network with two hidden layers:

```

data Network i h1 h2 o =
  Net { _nLayer1 :: !(Layer i h1)
       , _nLayer2 :: !(Layer h1 h2)
       , _nLayer3 :: !(Layer h2 o)
       }
  deriving (Show, Generic)

```

```

instance NFData (Network i h1 h2 o)
makeLenses ''Network

```

These are pretty straightforward container types...pretty much exactly the type you'd make to represent these networks! Note that, following true Haskell form, we separate out logic from data. This should be all we need.

## Instances

Things are much simpler if we had `Num` and `Fractional` instances for everything, so let's just go ahead and define that now, as well. Just a little bit of boilerplate.

```

instance (KnownNat i, KnownNat o) => Num (Layer i o) where
  Layer w1 b1 + Layer w2 b2 = Layer (w1 + w2) (b1 + b2)
  Layer w1 b1 - Layer w2 b2 = Layer (w1 - w2) (b1 - b2)
  Layer w1 b1 * Layer w2 b2 = Layer (w1 * w2) (b1 * b2)
  abs    (Layer w b)      = Layer (abs    w) (abs    b)
  signum (Layer w b)      = Layer (signum w) (signum b)
  negate (Layer w b)      = Layer (negate w) (negate b)
  fromInteger x           = Layer (fromInteger x) (fromInteger x)

```

```

instance (KnownNat i, KnownNat h1, KnownNat h2, KnownNat o) => Num (Network i h1 h2 o) where
  Net a b c + Net d e f = Net (a + d) (b + e) (c + f)
  Net a b c - Net d e f = Net (a - d) (b - e) (c - f)
  Net a b c * Net d e f = Net (a * d) (b * e) (c * f)
  abs (Net a b c)      = Net (abs a) (abs b) (abs c)
  signum (Net a b c)   = Net (signum a) (signum b) (signum c)
  negate (Net a b c)   = Net (negate a) (negate b) (negate c)
  fromInteger x        = Net (fromInteger x) (fromInteger x) (fromInteger x)

instance (KnownNat i, KnownNat o) => Fractional (Layer i o) where
  Layer w1 b1 / Layer w2 b2 = Layer (w1 / w2) (b1 / b2)
  recip (Layer w b)         = Layer (recip w) (recip b)
  fromRational x            = Layer (fromRational x) (fromRational x)

instance (KnownNat i, KnownNat h1, KnownNat h2, KnownNat o) => Fractional (Network i h1 h2 o) where
  Net a b c / Net d e f = Net (a / d) (b / e) (c / f)
  recip (Net a b c)      = Net (recip a) (recip b) (recip c)
  fromRational x         = Net (fromRational x) (fromRational x) (fromRational x)

```

KnownNat comes from *base*; it's a typeclass that *hmatrix* uses to refer to the numbers in its type and use it to go about its normal hmatrixy business.

## Ops

Now, *backprop* does require *primitive* differentiable operations on our relevant types to be defined. *backprop* uses these primitive operations to tie everything together. Ideally we'd import these from a library that implements these for you, and the end-user never has to make these primitives.

But in this case, I'm going to put the definitions here to show that there isn't any magic going on. If you're curious, refer to documentation for *Op* for more details on how *Op* is implemented and how this works.

First, matrix-vector multiplication primitive, giving an explicit gradient function.

```

infixr 8 #>!
(#>!)
  :: (KnownNat m, KnownNat n, Reifies s W)
  => BVar s (L m n)
  -> BVar s (R n)
  -> BVar s (R m)
(#>!) = liftOp2 . op2 $ \m v ->
  ( m #> v, \g -> (g `outer` v, tr m #> g) )

```

Dot products would be nice too.

```

infixr 8 <.>!
(<.>!)
  :: (KnownNat n, Reifies s W)
  => BVar s (R n)
  -> BVar s (R n)
  -> BVar s Double
(<.>!) = liftOp2 . op2 $ \x y ->
  ( x <.> y, \g -> (konst g * y, x * konst g)
  )

```

Also a function to fill a vector with the same element:

```

konst'
  :: (KnownNat n, Reifies s W)
  => BVar s Double
  -> BVar s (R n)
konst' = liftOp1 . op1 $ \c -> (konst c, HM.sumElements . extract)

```

Finally, an operation to sum all of the items in the vector.

```

sumElements'
  :: (KnownNat n, Reifies s W)
  => BVar s (R n)
  -> BVar s Double
sumElements' = liftOp1 . op1 $ \x -> (HM.sumElements (extract x), konst)

```

Here's the logistic function, which we'll use as an activation function for internal layers. But, because BVars have a Floating instance, we can just write it using typeclass functions.

```

logistic :: Floating a => a -> a
logistic x = 1 / (1 + exp (-x))
{-# INLINE logistic #-}

```

## Running our Network

Now that we have our primitives in place, let's actually write a function to run our network! And, once we do this, we automatically also have functions to back-propagate our network!

Normally, to write this function, we'd write:

```

runLayerNormal
  :: (KnownNat i, KnownNat o)
  => Layer i o
  -> R i
  -> R o
runLayerNormal l x = (l ^. lWeights) #> x + (l ^. lBiases)
{-# INLINE runLayerNormal #-}

```

Using the lWeights and lBiases lenses to access the weights and biases of our layer. However, we can translate this to *backprop* by operating on BVars instead of the type directly, and using our backprop-aware #>!:

```

runLayer
  :: (KnownNat i, KnownNat o, Reifies s W)
  => BVar s (Layer i o)
  -> BVar s (R i)
  -> BVar s (R o)
runLayer l x = (l ^^ lWeights) #>! x + (l ^^ lBiases)
{-# INLINE runLayer #-}

```

^. lets to access data within a value using a lens, and ^^ lets you access data within a BVar using a lens:

```

(^.) :: a -> Lens a b -> b
(^^.) :: BVar s a -> Lens a b -> BVar s b

```

Now runLayer is a function on two inputs that can be backpropagated, automatically! We can find its gradient given any input, and also run it to get our expected output as well.

Before writing our final network runner, we need a function to compute the “softmax” of our output vector. Writing it normally would look like:

```

softmaxNormal :: KnownNat n => R n -> R n
softmaxNormal x = konst (1 / HM.sumElements (extract x)) * exp x
{-# INLINE softmaxNormal #-}

```

But we can make the mechanical shift to the backpropagatable version:

```

softmax :: (KnownNat n, Reifies s W) => BVar s (R n) -> BVar s (R n)
softmax x = konst' (1 / sumElements' x) * exp x
{-# INLINE softmax #-}

```

With that in hand, let's compare how we would normally write a function to run our network:

```

runNetNormal
  :: (KnownNat i, KnownNat h1, KnownNat h2, KnownNat o)
  => Network i h1 h2 o
  -> R i
  -> R o
runNetNormal n = softmaxNormal
  . runLayerNormal (n ^. nLayer3)
  . logistic
  . runLayerNormal (n ^. nLayer2)
  . logistic
  . runLayerNormal (n ^. nLayer1)
{-# INLINE runNetNormal #-}

```

Basic function composition, neat. We use our lenses `nLayer1`, `nLayer2`, and `nLayer3` to extract the first, second, and third layers from our network.

Writing it in a way that backprop can use is also very similar:

```

runNetwork
  :: (KnownNat i, KnownNat h1, KnownNat h2, KnownNat o, Reifies s W)
  => BVar s (Network i h1 h2 o)
  -> BVar s (R i)
  -> BVar s (R o)
runNetwork n = softmax
  . runLayer (n ^^ nLayer3)
  . logistic
  . runLayer (n ^^ nLayer2)
  . logistic
  . runLayer (n ^^ nLayer1)
{-# INLINE runNetwork #-}

```

And now here again we use `^^`. (instead of `^.`) to extract a value from our `BVar` of a `Network`, using a lens.

## Computing Errors

Now, training a neural network is about calculating its gradient with respect to some error function. The library calculates the gradient for us – we just need to tell it how to compute the error function.

For classification problems, we usually use a cross entropy error. Given a target vector, how does our neural network's output differ from what is expected? Lower numbers are better!

Again, let's look at a “normal” implementation, regular variables and no backprop:

```

crossEntropyNormal :: KnownNat n => R n -> R n -> Double
crossEntropyNormal targ res = negate $ log res <.> targ
{-# INLINE crossEntropyNormal #-}

```

And we can see that the backpropable version is pretty similar. We see `constVar t`, to introduce a `BVar` that is a constant value (that we don't care about the gradient of).

```
crossEntropy
  :: (KnownNat n, Reifies s W)
  => R n
  -> BVar s (R n)
  -> BVar s Double
crossEntropy targ res = negate $ log res <.>! constVar targ
{-# INLINE crossEntropy #-}
```

Our final “error function”, then, is:

```
netErr
  :: (KnownNat i, KnownNat h1, KnownNat h2, KnownNat o, Reifies s W)
  => R o
  -> BVar s (Network i h1 h2 o)
  -> BVar s (R i)
  -> BVar s Double
netErr targ n = crossEntropy targ . runNetwork n
{-# INLINE netErr #-}
```

## The Magic

The actual “magic” of the library happens with the functions to “run” the functions we defined earlier:

```
evalBP   :: (forall s. Reifies s W => BVar s a -> BVar s b) -> a -> b
gradBP   :: (forall s. Reifies s W => BVar s a -> BVar s b) -> a -> a
backprop :: (forall s. Reifies s W => BVar s a -> BVar s b) -> a -> (b, a)
```

`evalBP` “runs” the function like normal, `gradBP` computes the gradient of the function, and `backprop` computes both the result and the gradient.

So, if we have a network `net0`, an input vector `x`, and a target vector `t`, we could compute its error using:

```
evalBP (uncurryVar netErr) (net0, x) :: Double
```

We use the `uncurryVar` helper function to translate our curried function into a function taking one tuple argument:

```
uncurryVar :: (BVar s a -> BVar s b -> BVar s c) -> (BVar s (a, b) -> BVar s c)
```

And we can calculate its *gradient* using:

```
gradBP (uncurryVar netErr) (net0, x) :: (Network i h1 h2 o, R i)
```

## Pulling it all together

Let's write a simple function to step our network in the direction opposite of the gradient to train our model:

```
trainStep
  :: forall i h1 h2 o. (KnownNat i, KnownNat h1, KnownNat h2, KnownNat o)
  => Double           -- ^ learning rate
  -> R i              -- ^ input
  -> R o              -- ^ target
  -> Network i h1 h2 o -- ^ initial network
  -> Network i h1 h2 o
```

```

trainStep r !x !t !n =
  case gradBP (uncurryVar (netErr t)) (n, x) of
    (gN, _) -> n - (realToFrac r * gN)
{-# INLINE trainStep #-}

```

Here's a convenient wrapper for training over all of the observations in a list:

```

trainList
  :: (KnownNat i, KnownNat h1, KnownNat h2, KnownNat o)
=> Double          -- ^ learning rate
-> [(R i, R o)]     -- ^ input and target pairs
-> Network i h1 h2 o -- ^ initial network
-> Network i h1 h2 o
trainList r = flip $ foldl' (\n (x,y) -> trainStep r x y n)

```

testNet will be a quick way to test our net by computing the percentage of correct guesses: (mostly using *hmatrix* stuff, so don't mind too much)

```

testNet
  :: forall i h1 h2 o. (KnownNat i, KnownNat h1, KnownNat h2, KnownNat o)
=> [(R i, R o)]
-> Network i h1 h2 o
-> Double
testNet xs n = sum (map (uncurry test) xs) / fromIntegral (length xs)
where
  test :: R i -> R o -> Double          -- test if the max index is correct
  test x (extract->t)
    | HM.maxIndex t == HM.maxIndex (extract r) = 1
    | otherwise                               = 0
  where
    r :: R o
    r = evalBP (uncurryVar runNetwork) (n , x)

```

And now, a main loop!

If you are following along at home, download the mnist data set files and uncompress them into the folder `data`, and everything should work fine.

```

main :: IO ()
main = MWC.withSystemRandom $ \g -> do
  Just train <- loadMNIST "data/train-images-idx3-ubyte" "data/train-labels-idx1-ubyte"
  Just test  <- loadMNIST "data/t10k-images-idx3-ubyte" "data/t10k-labels-idx1-ubyte"
  putStrLn "Loaded data."
  net0 <- MWC.uniformR @(Network 784 300 100 10) (-0.5, 0.5) g
  flip evalStateT net0 . forM_ [1..] $ \e -> do
    train' <- liftIO . fmap V.toList $ MWC.uniformShuffle (V.fromList train) g
    liftIO $ printf "[Epoch %d]\n" (e :: Int)

    forM_ ([1..] `zip` chunksOf batch train') $ \(b, chnk) -> StateT $ \n0 -> do
      printf "(Batch %d)\n" (b :: Int)

      t0 <- getCurrentTime
      n' <- evaluate . force $ trainList rate chnk n0
      t1 <- getCurrentTime
      printf "Trained on %d points in %.s.\n" batch (show (t1 `diffUTCTime` t0))

      let trainScore = testNet chnk n'

```

```

        testScore = testNet test n'
    printf "Training error:  %.2f%%\n" ((1 - trainScore) * 100)
    printf "Validation error: %.2f%%\n" ((1 - testScore ) * 100)

    return ((), n')
where
    rate = 0.02
    batch = 5000

```

Each iteration of the loop:

1. Shuffles the training set
2. Splits it into chunks of `batch` size
3. Uses `trainList` to train over the batch
4. Computes the score based on `testNet` based on the training set and the test set
5. Prints out the results

And, that's really it!

## Result

I haven't put much into optimizing the library yet, but the network (with hidden layer sizes 300 and 100) seems to take 25s on my computer to finish a batch of 5000 training points. It's slow (five minutes per 60000 point epoch), but it's a first unoptimized run and a proof of concept! It's my goal to get this down to a point where the result has the same performance characteristics as the actual backend (*hmatrix*), and so overhead is 0.

## Main takeaways

Most of the actual heavy lifting/logic actually came from the *hmatrix* library itself. We just created simple types to wrap up our bare matrices.

Basically, all that *backprop* did was give you an API to define *how to run* a neural net — how to *run* a net based on a `Network` and `R i` input you were given. The goal of the library is to let you write down how to run things in as natural way as possible.

And then, after things are run, we can just get the gradient and roll from there!

Because the heavy lifting is done by the data types themselves, we can presumably plug in *any* type and any tensor/numerical backend, and reap the benefits of those libraries' optimizations and parallelizations. *Any* type can be backpropagated! :D

## What now?

Check out the docs for the `Numeric.Backprop` module for a more detailed picture of what's going on, or find more examples at the github repo!

## Boring stuff

Here is a small wrapper function over the `mnist-idx` library loading the contents of the `idx` files into *hmatrix* vectors:



```

loadMNIST
  :: FilePath
  -> FilePath
  -> IO (Maybe [(R 784, R 10)])
loadMNIST fpI fpL = runMaybeT $ do
  i <- MaybeT $ decodeIDXFile      fpI
  l <- MaybeT $ decodeIDXLabelsFile fpL
  d <- MaybeT . return $ labeledIntData l i
  r <- MaybeT . return $ for d (bitraverse mkImage mkLabel . swap)
  liftIO . evaluate $ force r
where
  mkImage :: VU.Vector Int -> Maybe (R 784)
  mkImage = create . VG.convert . VG.map (\i -> fromIntegral i / 255)
  mkLabel :: Int -> Maybe (R 10)
  mkLabel n = create $ HM.build 10 (\i -> if round i == n then 1 else 0)

```

And here are instances to generating random vectors/matrices/layers/networks, used for the initialization step.

```

instance KnownNat n => MWC.Variate (R n) where
  uniform g = randomVector <$> MWC.uniform g <*> pure Uniform
  uniformR (l, h) g = (\x -> x * (h - l) + l) <$> MWC.uniform g

instance (KnownNat m, KnownNat n) => MWC.Variate (L m n) where
  uniform g = uniformSample <$> MWC.uniform g <*> pure 0 <*> pure 1
  uniformR (l, h) g = (\x -> x * (h - l) + l) <$> MWC.uniform g

instance (KnownNat i, KnownNat o) => MWC.Variate (Layer i o) where
  uniform g = Layer <$> MWC.uniform g <*> MWC.uniform g
  uniformR (l, h) g = (\x -> x * (h - l) + l) <$> MWC.uniform g

instance (KnownNat i, KnownNat h1, KnownNat h2, KnownNat o) => MWC.Variate (Network i h1 h2 o) where
  uniform g = Net <$> MWC.uniform g <*> MWC.uniform g <*> MWC.uniform g
  uniformR (l, h) g = (\x -> x * (h - l) + l) <$> MWC.uniform g

```

Also for now we do need an orphan Num instance on tuples:

```

instance (Num a, Num b) => Num (a, b) where
  (x1,y1) + (x2,y2) = (x1 + x2, y1 + y2)
  (x1,y1) * (x2,y2) = (x1 * x2, y1 * y2)
  (x1,y1) - (x2,y2) = (x1 - x2, y1 - y2)
  abs (x, y)       = (abs x, abs y)
  signum (x, y)     = (signum x, signum y)
  fromInteger x     = (fromInteger x, fromInteger x)

```